

# Reinforcement Learning for Texas Hold'em Poker

Michele Sanna

# Motivation and Focus

This project explores the application of reinforcement learning to Heads-Up Texas Hold'em Poker, a zero-sum, imperfect information game. The focus lies on algorithmic performance rather than network architecture, using no human data.

# Simplifying the Game

To handle complexity:

- Betting is discretized.
- Only the two-player (Heads-Up) variant is considered.
- Card information is abstracted using hand equity — the estimated probability of winning.

This abstraction avoids sparse card representations and accelerates learning.

The observation space of the agents for this work will consist in:

- A one-hot encoded representation of the game phase (i.e. how many community cards are revealed)
- The equity of the hand of the agent
- The normalized size of the pot
- The normalized size of both players stack
- The bet of the player
- The bet of the opponent
- A condensed betting history of the opponent (the total opponent bet for each game phase)

# Equity Estimation

Monte Carlo simulations (2k–8k samples) are used to compute hand equity. Though computationally expensive, this significantly improves training convergence. Accuracy compared to online calculators shows a mean deviation of just 3.9%.

# DQN: Deep Q-Network

**DQN** uses a neural network to estimate action values:

$$Q(s, a; \theta) \approx Q^*(s, a)$$

During training, the target is computed as:

$$y = r + \gamma \max_{a'} Q(s', a'; \theta)$$

**Problem:** The max operator can select and evaluate the same overestimated value, leading to:

- Overoptimistic Q-values
- Instability in training
- Poor performance in stochastic or noisy environments

# Double DQN: Reducing Overestimation Bias

**Double DQN** decouples action selection and evaluation:

$$y = r + \gamma Q(s', \arg \max_{a'} Q(s', a'; \theta), \theta^-)$$

Where:


- $\theta$  is the current network (used for action selection)
- $\theta^-$  is the target network (used for action evaluation)

**Benefits:**

- Reduces overestimation bias
- Improves stability
- Often leads to better policies in practice

# Naive Approach: DDQN vs DDQN

Two Double DQN agents were trained by self-play. Despite variations in hyperparameters, all learned a suboptimal strategy: never folding. This reflects a known limitation—single-agent RL algorithms often fail in multi-agent, imperfect-information games.



**Figure:** Folding behavior during naive DDQN vs DDQN training



# Limitations of Standard RL

In imperfect-information settings:

- Optimal strategies are often stochastic.
- Best-response policies can be exploitable.
- Classical RL algorithms lack convergence guarantees to Nash equilibria.

These factors cause instability or cycling in self-play dynamics.

# Better Solutions: NFSP over ReBel and MMD

Several approaches address imperfect-information games in deep RL:

**ReBel:** Combines reinforcement learning with search and belief state modeling.

**MMD:** Uses a modified PPO with tailored regularization.

**NFSP:** Based on *fictitious play*, where agents respond to the opponent's average strategy.

**This work uses NFSP** due to its:

- Simpler implementation than ReBel
- Simpler math than MMD

# A More Principled Approach: NFSP

Neural Fictitious Self-Play (NFSP) adapts fictitious play to deep RL. It mixes:

- A best-response DQN policy (used with probability  $\eta$ )
- An average-policy network (used with probability  $1 - \eta$ )

Agents alternate between these strategies, gradually converging to more stable behavior.

# NFSP Algorithm Overview

---

**Algorithm 1** Neural Fictitious Self-Play (NFSP) with fitted Q-learning

---

Initialize game  $\Gamma$  and execute an agent via RUNAGENT for each player in the game

**function** RUNAGENT( $\Gamma$ )

    Initialize replay memories  $\mathcal{M}_{RL}$  (circular buffer) and  $\mathcal{M}_{SL}$  (reservoir)

    Initialize average-policy network  $\Pi(s, a | \theta^\Pi)$  with random parameters  $\theta^\Pi$

    Initialize action-value network  $Q(s, a | \theta^Q)$  with random parameters  $\theta^Q$

    Initialize target network parameters  $\theta^{Q'} \leftarrow \theta^Q$

    Initialize anticipatory parameter  $\eta$

**for each** episode **do**

        Set policy  $\sigma \leftarrow \begin{cases} \epsilon\text{-greedy}(Q), & \text{with probability } \eta \\ \Pi, & \text{with probability } 1 - \eta \end{cases}$

        Observe initial information state  $s_1$  and reward  $r_1$

**for**  $t = 1, T$  **do**

            Sample action  $a_t$  from policy  $\sigma$

            Execute action  $a_t$  in game and observe reward  $r_{t+1}$  and next information state  $s_{t+1}$

            Store transition  $(s_t, a_t, r_{t+1}, s_{t+1})$  in reinforcement learning memory  $\mathcal{M}_{RL}$

**if** agent follows best response policy  $\sigma = \epsilon\text{-greedy}(Q)$  **then**

                Store behaviour tuple  $(s_t, a_t)$  in supervised learning memory  $\mathcal{M}_{SL}$

**end if**

            Update  $\theta^\Pi$  with stochastic gradient descent on loss

$\mathcal{L}(\theta^\Pi) = \mathbb{E}_{(s,a) \sim \mathcal{M}_{RL}} [-\log \Pi(s, a | \theta^\Pi)]$

            Update  $\theta^Q$  with stochastic gradient descent on loss

$\mathcal{L}(\theta^Q) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{M}_{RL}} \left[ \left( r + \max_{a'} Q(s', a' | \theta^{Q'}) - Q(s, a | \theta^Q) \right)^2 \right]$

            Periodically update target network parameters  $\theta^{Q'} \leftarrow \theta^Q$

**end for**

**end for**

**end function**

---

Figure: Original NFSP algorithm

This implementation uses DDQN as the underlying RL method, and equity-based state representation.

# Training Comparison

Both NFSP and DDQN-vs-DDQN agents were trained for 150k episodes. The DDQN agents collapsed into non-folding behavior, while NFSP maintained a balanced strategy.





Figure: Fold frequency over time

# Approximate Exploitability

Exploitability was approximated by training a DDQN best-response.




Figure: Approximate exploitability

# Performance vs Randomized Opponents

Each agent was evaluated against:

- A uniformly random player
- A player that always goes all-in
- A conservative bettor (check/small bet)

# Performance vs Randomized Opponents




Figure: Performance vs random strategy



# Performance vs Randomized Opponents




Figure: Performance vs all-in strategy

# Performance vs Randomized Opponents




Figure: Performance vs conservative better

# Head-to-Head

Episodi	mBb per hand
0	0.0
10000	-1216.2
20000	-1645.1
30000	768.475
40000	-10662.34
50000	123.68
60000	1355.93
70000	369.98
80000	472.78
90000	200.22
100000	135.68
110000	933.87
120000	-67.15
130000	615.74
140000	1469.11
150000	449.6

**Figure:** DDQN vs. DDQN advantage over NFSP. Head-to-head performance over training